

RAJ

☎ +91 6387287539 | ✉ rajr59110@gmail.com | 🔗 LinkedIn | 🐙 GitHub | 🌐 Portfolio

PROFESSIONAL SUMMARY

Data Analyst specializing in enterprise data governance, modern data stack implementation (GCP, AWS, Snowflake), and full-cycle ELT automation. Proven ability to architect scalable data lakehouses and hybrid pipelines (dbt Core, DuckDB, Parquet) to efficiently process and validate massive datasets (400M+ rows). Highly capable technical executor focused on establishing robust data standardization, strict quality frameworks, and secure data products that empower enterprise stakeholders to make high-stakes, compliant business decisions.

EXPERIENCE

Opsie

DATA ANALYST

May 2025 – March 2026

Internship : Hybrid , Mohali , Punjab

- Orchestrated scalable ELT workflows to automate the secure ingestion and standardization of complex, multi-source enterprise datasets into Google Cloud Storage (GCS) for centralized data governance.
- Engineered optimized data architectures in Google BigQuery, utilizing advanced SQL to materialize external data into high-performance native tables, effectively managing critical enterprise data assets.
- Designed and deployed modular, version-controlled transformation pipelines using dbt Core, establishing robust data quality processes, strict testing, and complete lineage tracking.
- Deployed production-ready, governed data products seamlessly integrated with downstream platforms, eliminating manual overhead and enabling secure access for cross-functional end-users.

PROJECTS

Promotional Waste Audit & Compute Optimization | GCP, BigQuery, DuckDB, dbt, Power BI

- * Architected an automated ELT pipeline processing 411M+ rows of complex, multi-table transactional data, implementing strict data validation rules to detect anomalies and ensure enterprise-level data integrity.
- * Engineered a scalable data lake in Google Cloud Storage utilizing Parquet format for optimized columnar reading, and pioneered a hybrid SQL/Python exploratory workflow with DuckDB in Colab Enterprise to drastically reduce BigQuery compute costs.
- * Automated complex data transformations and table staging (External to Native) using dbt Core and VS Code, establishing robust data quality testing and version-controlled governance frameworks.
- * Deployed secure, dynamic Power BI data products utilizing advanced DAX to monitor data anomalies, enforce standardization, and equip stakeholders with highly accurate reporting assets.

SaaS LTV Optimization & Optimizing Marketing Spend | AWS S3, Athena, DuckDB, Parquet, dbt, Power BI

- Engineered a scalable relational data model analyzing 22M+ transaction records and 6M member profiles for a B2C SaaS platform, establishing strict data standardization protocols to ensure accurate downstream analytics.
- Architected a serverless, decoupled data lakehouse by storing heavy transactional datasets in optimized Parquet format within Amazon S3, querying directly using Amazon Athena to minimize compute overhead.
- Pioneered a hybrid DuckDB workflow (SQL + Python) within VS Code to perform lightning-fast, in-memory exploratory data analysis on the 22M row dataset, radically cutting down interactive query latency.
- Automated production-grade transformations and standardized user retention logic by building modular dbt Core pipelines, ensuring reliable integration with the AWS Athena catalog.

Market Basket Analysis & Cross-Selling Engine | Snowflake, AWS S3, Tableau, dbt, DuckDB

- Engineered the data architecture for a cross-selling recommendation engine, performing Market Basket Analysis on 2M+ transactions and ensuring high-fidelity data staging for downstream modeling.
- Architected a scalable pipeline by orchestrating Parquet datasets into an AWS S3 data lake and staging them into Snowflake via highly optimized external and native tables.

- Pioneered a hybrid exploratory workflow utilizing DuckDB and Snowflake Notebooks (SQL/Python) to rapidly analyze entity combinations and compute conditional probabilities at scale.
- Automated complex association rule logic using dbt Core and VS Code, establishing strict version control and reliable data transformations to support critical business strategies.

Competitive Intelligence & Market Share Analytics | *DuckDB, Python, SQL, Parquet, Power BI*

- Spearheaded data processing for a competitive intelligence initiative, analyzing 10M+ rows of complex retail data to model market dynamics and establish reliable, standardized baseline metrics.
- Engineered a high-performance, zero-infrastructure data pipeline within Google Colab, leveraging Parquet files and a DuckDB hybrid SQL/Python workflow to process massive datasets entirely in-memory.
- Executed geospatial data modeling across target counties, standardizing location-based datasets to ensure high accuracy for spatial analytics and distribution mapping.
- Diagnosed store-level data anomalies and velocity gaps, deploying standardized, localized data models that equipped cross-functional teams with secure, high-fidelity intelligence.

EDUCATION

Chandigarh University
Bachelor of Engineering

Sept. 2021 – June 2025
Mohali, Punjab

SKILLS

Languages & Querying: Python (Pandas), SQL, Microsoft Excel, DAX

Cloud Platforms: Google Cloud Platform (BigQuery, Cloud Storage), AWS (S3, Athena), Snowflake

Data Engineering Architecture: dbt (Data Build Tool), dbt Core, DuckDB (Hybrid), Git/GitHub

Visualization & Reporting: Power BI, Tableau, Seaborn, Matplotlib

Data Notebooks: Google Colab Enterprise, Snowflake Notebooks, Amazon SageMaker

CERTIFICATIONS

Oracle Cloud Infrastructure 2025 Certified Data Science Professional | *Oracle*

Google Cloud Data Analytics Certificate | *Google*