RAJ

+91 6387287539 | rajr59110@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Data Analyst specializing in modern data stack implementation (GCP, AWS, Snowflake) and full-cycle ELT automation. Proven ability to engineer scalable data lakehouses and hybrid in-memory pipelines (DuckDB, Parquet, dbt) to efficiently process massive datasets (400M+ rows). Highly capable technical executor focused on translating complex data into actionable business intelligence, developing advanced behavioral segmentations and analytical frameworks (LTV, Market Basket Analysis) that equip stakeholders to optimize marketing ROI and improve profit margins.

EXPERIENCE

Opsie — **July 2025 – June 2026**
DATA ANALYST — *Remote*

- Orchestrated scalable ELT (Extract, Load, Transform) workflows to automate the ingestion of complex, multi-source enterprise datasets into Google Cloud Storage (GCS) for centralized data staging.
- Engineered optimized data architectures in Google BigQuery, utilizing advanced SQL to materialize external data into high-performance native tables, significantly reducing query latency.
- Designed and deployed modular, version-controlled transformation pipelines using dbt Core, ensuring strict data testing, lineage tracking, and reliability.
- Deployed production-ready data marts seamlessly integrated with business intelligence (BI) platforms, eliminating manual overhead and automating real-time analytics reporting.

PROJECTS

Promotional Waste Audit & Margin Optimization | *GCP, BigQuery, DuckDB, dbt, Power BI*

* Architected an automated ELT pipeline processing 411M+ rows with multi-table of e-commerce behavioral data, diagnosing a $2M cart abandonment bleed and optimizing a margin-eroding discount strategy.
* Engineered a scalable data lake in Google Cloud Storage utilizing Parquet format for optimized columnar reading, and pioneered a hybrid SQL/Python exploratory workflow with DuckDB in Colab Enterprise to drastically reduce BigQuery compute costs.
* Automated complex data transformations and table staging (External to Native) using dbt Core and VS Code, ensuring highly reliable and version-controlled data pipelines.
* Deployed dynamic Power BI dashboards utilizing advanced DAX to suppress automatic 10% discounts for Safe Buyers and halt paid retargeting for Window Shoppers, significantly improving ROAS.

SaaS LTV Optimization & Optimizing Marketing Spend | *AWS S3, Athena, DuckDB, Parquet, dbt, Power BI*

- Engineered an LTV-based acquisition model analyzing 22M+ transaction records and 6M members detail with multi-table for a B2C SaaS platform, identifying severe marketing cash burn and strategic reallocation metrics toward high-retention "Power Users".
- Architected a serverless, decoupled data lakehouse by storing heavy user ledgers in optimized Parquet format within Amazon S3 and querying datasets directly using Amazon Athena to minimize compute overhead
- Pioneered a hybrid DuckDB workflow (SQL + Python) within VS Code to perform lightning-fast, in-memory exploratory data analysis on the 22M row dataset, radically cutting down interactive query latency.
- Automated production-grade transformations and 30-day strict subscriber churn logic by building modular dbt Core pipelines seamlessly integrated with the Athena catalog

Market Basket Analysis & Cross-Selling Engine | *Snowflake, AWS S3, Tableau, dbt, DuckDB*

- Engineered a cross-selling recommendation engine by performing Market Basket Analysis on 2M+ transactions and 92K products to maximize Average Order Value (AOV).
- Architected a scalable pipeline by dumping Parquet datasets into an AWS S3 data lake and staging them into Snowflake via highly optimized external and native tables.

- Pioneered a hybrid exploratory workflow utilizing DuckDB and Snowflake Notebooks (SQL/Python) to rapidly analyze item combinations and compute conditional probabilities for "Directional Pull" between lead and follow products.
- Automated association rule logic using dbt Core and VS Code to differentiate between low-retention impulse buys and habitual reorders, directly informing high-value "Subscribe & Save" business strategies.

Competitive Intelligence & Market Share Analytics | *DuckDB, Python, SQL, Parquet, Power BI*

- Spearheaded a competitive intelligence initiative analyzing 10M+ rows of retail data to pinpoint revenue bleed, market share erosion, and physical shelf-space loss against primary competitors.
- Engineered a high-performance, zero-infrastructure data pipeline within Google Colab, leveraging Parquet files and a DuckDB hybrid SQL/Python workflow to process massive datasets entirely in-memory.
- Executed geospatial distribution gap analysis across target counties to identify specific cities and retail whitespace where competitor products possessed wider availability.
- Diagnosed store-level product assortment and sales velocity gaps, equipping field sales representatives with actionable, localized intelligence to reclaim lost retail accounts and drive immediate revenue.

EDUCATION

Chandigarh University — Sept. 2021 – June 2025
Bachelor of Engineering — Mohali, Punjab

SKILLS

Languages & Querying: Python (Pandas), SQL (Advanced), DAX
Cloud Platforms: Google Cloud Platform (BigQuery, Cloud Storage), AWS (S3, Athena), Snowflake
Data Engineering Architecture: dbt (Data Build Tool), dbt Core, DuckDB (Hybrid), Git/GitHub
Visualization & Reporting: Power BI, Tableau, Seaborn, Matplotlib
Data Notebooks: Google Colab Enterprise, Snowflake Notebooks, Amazon SageMaker

CERTIFICATIONS

Oracle Cloud Infrastructure 2025 Certified Data Science Professional | *Oracle*
Google Cloud Data Analytics Certificate | *Google*